

# mark\_goldstein\_mg3479\_A1\_code

February 13, 2020

```
[1]: # Import dependencies
import random
import numpy as np
import torch
import torch.nn as nn

# You can find Alfredo's plotting code in plot_lib.py in this directory .
# Download it along with this assignment and keep it in the same directory.
from plot_lib import set_default, show_scatterplot, plot_bases

from matplotlib.pyplot import plot, title, axis

[2]: set_default()

[3]: # Set up your device
cuda = torch.cuda.is_available()
device = torch.device("cuda:0" if cuda else "cpu")

[4]: # Set up random seed to 1008. Do not change the random seed.
# Yes, these are all necessary when you run experiments!
seed = 1008
random.seed(seed)
np.random.seed(seed)
torch.manual_seed(seed)
if cuda:
    torch.cuda.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)
    torch.backends.cudnn.benchmark = False
    torch.backends.cudnn.deterministic = True
```

## 1 1. Full, slice, fill

Write a function `warm_up` that returns the 2D tensor with integers below. **Do not use any loops.**

```
1 2 1 1 1 1 2 1 1 1 1 2 1
2 2 2 2 2 2 2 2 2 2 2 2 2
1 2 1 1 1 1 2 1 1 1 1 2 1
```

```

1 2 1 3 3 1 2 1 3 3 1 2 1
1 2 1 3 3 1 2 1 3 3 1 2 1
1 2 1 1 1 1 2 1 1 1 1 2 1
2 2 2 2 2 2 2 2 2 2 2 2 2
1 2 1 1 1 1 2 1 1 1 1 2 1
1 2 1 3 3 1 2 1 3 3 1 2 1
1 2 1 3 3 1 2 1 3 3 1 2 1
1 2 1 1 1 1 2 1 1 1 1 2 1
2 2 2 2 2 2 2 2 2 2 2 2 2
1 2 1 1 1 1 2 1 1 1 1 2 1

```

Hint: Use `torch.full`, `torch.fill_`, and the slicing operator.

```

[5]: def warm_up():
    #raise NotImplementedError()
    t = torch.full((13, 13), 1)
    torch.fill_(t[1], 2)
    torch.fill_(t[11], 2)
    torch.fill_(t[6], 2)
    torch.fill_(t[:,1], 2)
    torch.fill_(t[:,11], 2)
    torch.fill_(t[:,6], 2)
    torch.fill_(t[3:5, 3:5], 3)
    torch.fill_(t[3:5, 8:10], 3)
    torch.fill_(t[8:10, 3:5], 3)
    torch.fill_(t[8:10, 8:10], 3)
    return t

    # Uncomment line below once you implement this function.
    print(warm_up())

```

```

tensor([[1., 2., 1., 1., 1., 1., 2., 1., 1., 1., 1., 2., 1.],
        [2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2.],
        [1., 2., 1., 1., 1., 1., 2., 1., 1., 1., 1., 2., 1.],
        [1., 2., 1., 3., 3., 1., 2., 1., 3., 3., 1., 2., 1.],
        [1., 2., 1., 3., 3., 1., 2., 1., 3., 3., 1., 2., 1.],
        [1., 2., 1., 1., 1., 1., 2., 1., 1., 1., 1., 2., 1.],
        [2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2.],
        [1., 2., 1., 1., 1., 1., 2., 1., 1., 1., 1., 2., 1.],
        [1., 2., 1., 3., 3., 1., 2., 1., 3., 3., 1., 2., 1.],
        [1., 2., 1., 3., 3., 1., 2., 1., 3., 3., 1., 2., 1.],
        [1., 2., 1., 1., 1., 1., 2., 1., 1., 1., 1., 2., 1.],
        [2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2.],
        [1., 2., 1., 1., 1., 1., 2., 1., 1., 1., 1., 2., 1.]])

```

## 2 2. To Loop or not to loop

The motivation for the following three sub-questions is to get you to think critically about how to write your deep learning code. These sorts of choices can make the difference between tractable and intractable model training.

### 2.1 2.1. mul\_row\_loop

Write a function `mul_row_loop`, using python loops with simple indexing but no advanced indexing/slicing, that receives a 2D tensor as input and returns a tensor of same size that is - equal to the input on the first row - 2 times the input's second row on the second row - 3 times the input's third row on the third row - etc..

For instance:

```
>>> t = torch.full((4, 8), 2.0)
>>> t
tensor([[2., 2., 2., 2., 2., 2., 2., 2.],
        [2., 2., 2., 2., 2., 2., 2., 2.],
        [2., 2., 2., 2., 2., 2., 2., 2.],
        [2., 2., 2., 2., 2., 2., 2., 2.]])
>>> mul_row(t)
tensor([[2., 2., 2., 2., 2., 2., 2., 2.],
        [4., 4., 4., 4., 4., 4., 4., 4.],
        [6., 6., 6., 6., 6., 6., 6., 6.],
        [8., 8., 8., 8., 8., 8., 8., 8.]])
```

```
[6]: def mul_row_loop(input_tensor):
      #raise NotImplementedError()
      n = len(input_tensor)
      m = len(input_tensor[0])
      ans = torch.zeros([n, m])
      for i in range(n):
          ans[i] = input_tensor[i] * (i + 1)
      return ans
```

```
[7]: t = torch.full((4, 8), 2.0)
      mul_row_loop(t)
```

```
[7]: tensor([[2., 2., 2., 2., 2., 2., 2., 2.],
          [4., 4., 4., 4., 4., 4., 4., 4.],
          [6., 6., 6., 6., 6., 6., 6., 6.],
          [8., 8., 8., 8., 8., 8., 8., 8.]])
```

### 3 2.2. mul\_row\_fast

Write a second version of the same function named `mul_row_fast` which uses tensor operations and no looping.

**Hint:** Use broadcasting and `torch.arange`, `torch.view`, and `torch.mul`.

```
[8]: def mul_row_fast(input_tensor):
      #raise NotImplementedError()
      multiplier = torch.arange(1, len(input_tensor) + 1).view(len(input_tensor), 1).float()
      return torch.mul(input_tensor, multiplier)
```

```
[9]: mul_row_fast(t)
```

```
[9]: tensor([[2., 2., 2., 2., 2., 2., 2., 2.],
            [4., 4., 4., 4., 4., 4., 4., 4.],
            [6., 6., 6., 6., 6., 6., 6., 6.],
            [8., 8., 8., 8., 8., 8., 8., 8.]])
```

## 4 2.3. times

Write a function `times` which takes a 2D tensor as input and returns the run times of `mul_row_loop` and `mul_row_fast` on this tensor, respectively. Use `time.perf_counter`.

Use `torch.ones` to create a 2D tensor of size (1000, 400) full of ones and run `times` on it (there should be more than two orders of magnitude difference).

```
[10]: from time import perf_counter
      def times(input_tensor):
          #raise NotImplementedError()
          start1 = perf_counter()
          mul_row_loop(input_tensor)
          end1 = perf_counter()
          t1 = end1 - start1
          start2 = perf_counter()
          mul_row_fast(input_tensor)
          end2 = perf_counter()
          t2 = end2 - start2
          return t1, t2

      # Uncomment lines below once you implement this function.
      input_tensor = torch.ones([1000, 400])
      time_1, time_2 = times(input_tensor)
      print('{} , {}'.format(time_1, time_2))
```

```
0.016407604000000298, 0.00042330800000023316
```

## 5 3. Non-linearities

In this section, we explore similar concepts to Lab 1 and get comfortable initializing modules like `nn.Linear` and using non-linearities in PyTorch.

## 5.1 3.1. ReLU

ReLU (Rectified Linear Unit) is a non-linear activation function defined as:

$$y = \max(0, x)$$

Define a fully connected neural network `linear_fc_relu` which: - takes 2 dimensional data as input and passes it through linear modules (`torch.nn.Linear`) - has one hidden layer of dimension 5 - has output dimension of 2 - has ReLU as an activation function

Create a tensor with input data `X` of size (100, 2) using `torch.randn`.


Following the example in [https://github.com/Atcold/pytorch-Deep-Learning-Minicourse/blob/master/02-space\\_stretching.ipynb](https://github.com/Atcold/pytorch-Deep-Learning-Minicourse/blob/master/02-space_stretching.ipynb), visualize the output of passing `X` through the neural network `linear_fc_relu`.

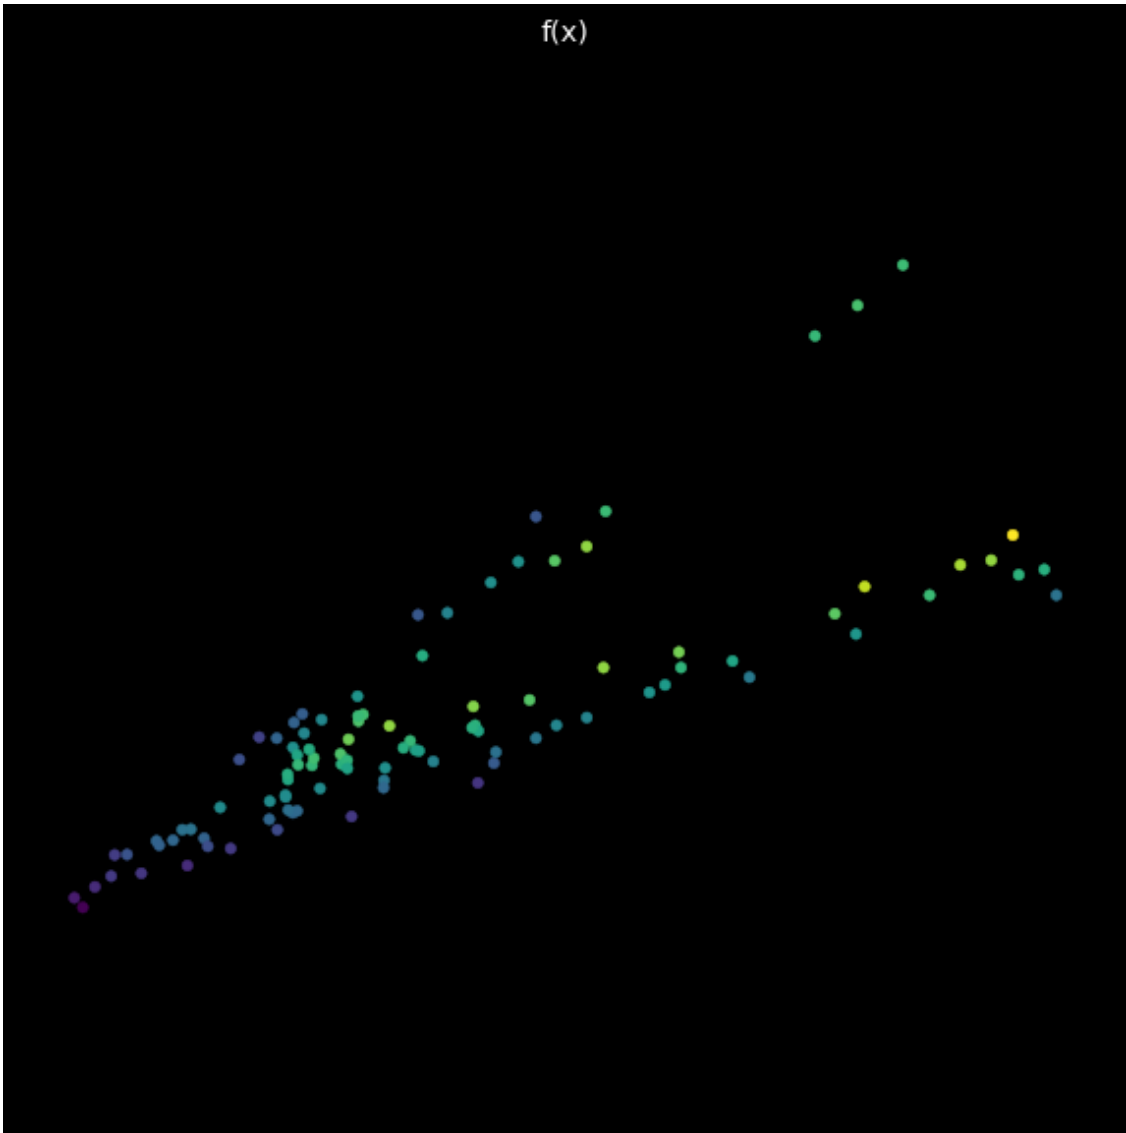
You can find Alfredo's plotting code in `plot_lib.py` in this directory. Download it along with this assignment and keep it in the same directory.


```
[11]: # Input data
X = torch.randn([100, 2])

[12]: # create 1-layer neural networks with ReLU activation
linear_fc_relu = nn.Sequential(
    nn.Linear(2, 5),
    nn.ReLU(),
    nn.Linear(5, 2)
)
# Visualize: TODO


[13]: colors = X[:, 0]
show_scatterplot(X, colors, title='x')
n_hidden = 5
NL = nn.ReLU()
for i in range(5):
    model = nn.Sequential(
        nn.Linear(2, n_hidden),
        NL,
        nn.Linear(n_hidden, 2)
    )
    model.to(device)
    with torch.no_grad():
        Y = model(X)
    show_scatterplot(Y, colors, title='f(x)')
```

















## 5.2 3.2. Sigmoid

The sigmoid function is another popular choice for a non-linear activation function which maps its input to values in the interval  $(0, 1)$ . It is formally defined as:

$$\sigma(x) = \frac{1}{1 + \exp[-x]}$$

Define a new neural network `linear_fc_sigmoid` which is the same architecture as in part 3.1. but with a sigmoid unit instead of ReLU.


Using the same `X` as in part 3.1, visualize the output of passing `X` through the neural network `linear_fc_sigmoid`.

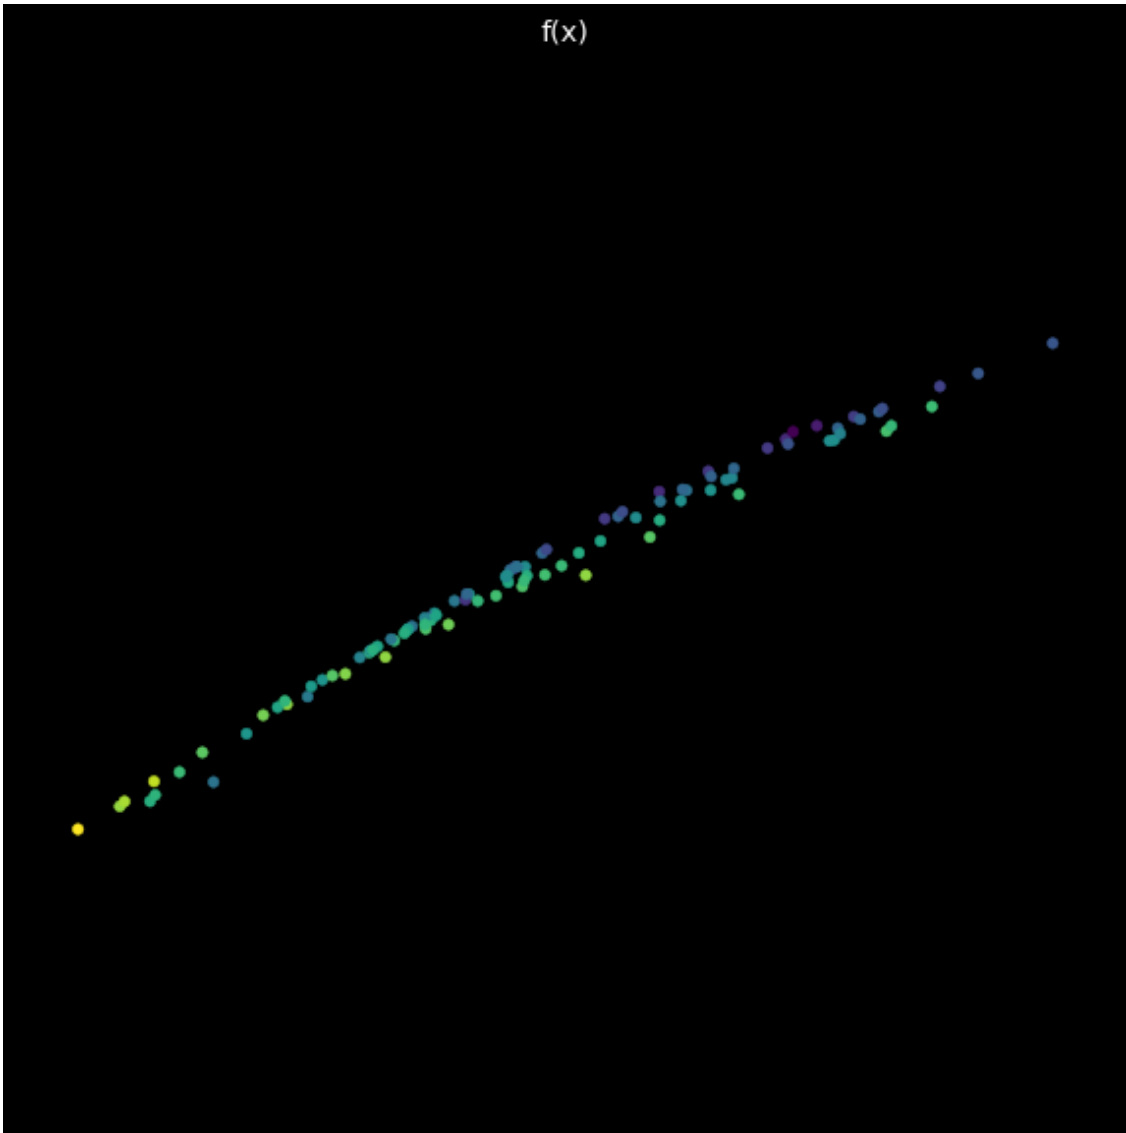
[14]: *# create 1-layer neural networks with Sigmoid activation*


```
linear_fc_sigmoid = nn.Sequential(  
    nn.Linear(2, 5),  
    nn.Sigmoid(),  
    nn.Linear(5, 2)  
  
)  
# Visualize: TODO
```


[15]:

```
colors = X[:, 0]  
show_scatterplot(X, colors, title='x')  
n_hidden = 5  
NL = nn.Sigmoid()  
  
for i in range(5):  
    model = nn.Sequential(  
        nn.Linear(2, n_hidden),  
        NL,  
        nn.Linear(n_hidden, 2)  
    )  
    model.to(device)  
    with torch.no_grad():  
        Y = model(X)  
    show_scatterplot(Y, colors, title='f(x)')
```

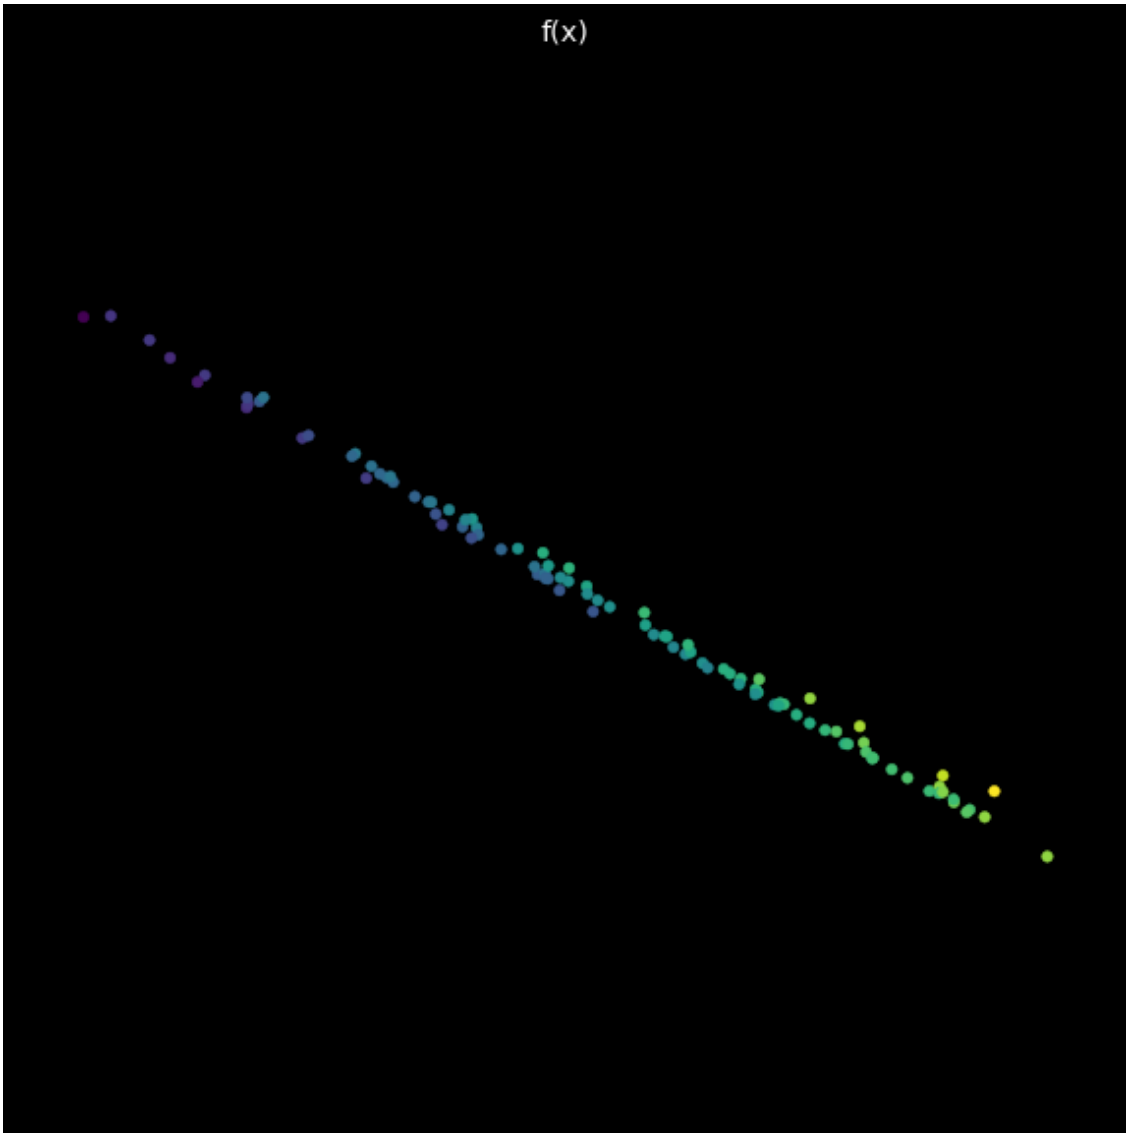















`[]:`